

Week 5 Demonstration

Tidy & Manipulate: Part II - Manipulate

Dr. Anil Dolgun

28/03/2018

1 / 30

Course Announcements

- Marking and feedback for Assignment 1 will be finalised by the next week.

Course Announcements

- Marking and feedback for Assignment 1 will be finalised by the next week.
- Suggested solutions to Week 4 Worksheet are available [here](#)

Course Announcements

- Marking and feedback for Assignment 1 will be finalised by the next week.
- Suggested solutions to Week 4 Worksheet are available [here](#)
- Details for Assignment 2 will be released at the end of this week.

Course Announcements

- Marking and feedback for Assignment 1 will be finalised by the next week.
- Suggested solutions to Week 4 Worksheet are available [here](#)
- Details for Assignment 2 will be released at the end of this week.
- Any other questions or concerns?


Tidy & Manipulate: Part II - Manipulate

3 / 30

Recall: Tidy Data Principles

Previous week:


- The framework of "tidy data" provides a standard and consistent way of storing data that makes transformation, visualization, and modeling easier.



Recall: Tidy Data Principles

Previous week:

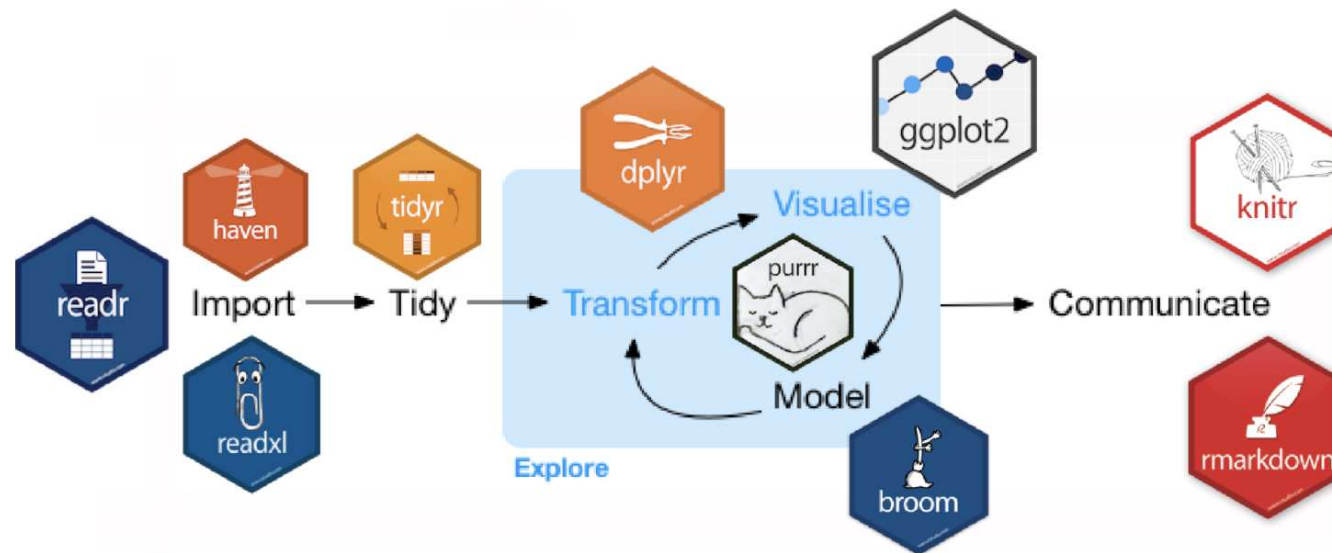
- The framework of "tidy data" provides a standard and consistent way of storing data that makes transformation, visualization, and modeling easier.



- Each variable must have its own column.
- Each observation must have its own row.
- Each value must have its own cell.

Recall: tidy

- `tidy` is a one such package which was built for the sole purpose of simplifying the process of creating tidy data.
- Following tidy principles makes manipulation, transformation, visualization, and modeling easier.
- `tidyverse` is a set of packages that work in harmony because they share common data representations and API design



The grammar of Data Manipulation: dplyr

- There are many data manipulation packages/functions in R,
- Most of them lack consistent coding and the ability to easily flow together.
- This leads to difficult-to-read nested functions and/or choppy code.

The grammar of Data Manipulation: dplyr

- There are many data manipulation packages/functions in R,
- Most of them lack consistent coding and the ability to easily flow together.
- This leads to difficult-to-read nested functions and/or choppy code.
- The `dplyr` package is regarded as the "tidy" grammar in R.
- It provides a consistent set of verbs that help you solve the most common data manipulation challenges.
- Remember: `dplyr` functions work with pipes `%>%` and expect a data frame as input.

The grammar of Data Manipulation: dplyr

- There are six fundamental functions of data manipulation that `dplyr` provides:
 - `select()` pick/select variables
 - `filter()` pick/filter observations based on values
 - `arrange()` sort variables
 - `mutate()` create new variables
 - `summarise()` summarise data by functions of choice
 - `group_by()` + `summarise()`
- There are also functions to join and merge data sets:
 -
 -
 -
 -
- The full list of capabilities can be found in the [dplyr reference manual](#).
- I highly recommend going through it as there are many great functions provided by `dplyr` that I will not cover here.

select() : select variables

- Often we only assess specific variables. The `select()` function allows us to select variables.
- In addition to the existing functions like `:` and `c()`, there are a number of special functions that can work inside select.

<code>-</code>	Select everything but
<code>:</code>	Select range
<code>contains()</code>	Select columns whose name contains a character string
<code>ends_with()</code>	Select columns whose name ends with a string
<code>everything()</code>	Select every column
<code>matches()</code>	Select columns whose name matches a regular expression
<code>num_range()</code>	Select columns named x1, x2, x3, x4, x5
<code>one_of()</code>	Select columns whose names are in a group of names
<code>starts_with()</code>	Select columns whose name starts with a character string

Class Activity:

- Visit <https://b.socrative.com/login/student/>
- Room Name: MATH2349



- Enter a nickname
- Work in small groups
- Import the [CustomerData.csv](#) data set and answer the questions on Socrative.
- No need to type any commands in R!

Import Customer Data

- The `CustomerData.csv` data set includes some characteristics of 5000 customers. Header of this data set is as follows.

```
customer <- read_csv("../data/CustomerData.csv")
head(customer[, 1:7], 5)
```

```
## # A tibble: 5 x 7
##   CustomerID      Region TownSize Gender   Age EducationYears JobCategory
##   <chr>          <int> <chr>   <chr> <int>         <int> <chr>
## 1 3964-QJWTRG-NPN      1 2      Female   20             15 Professional
## 2 0648-AIPJSP-UVM      5 5      Male     22             17 Sales
## 3 5195-TLUDJE-HVO      3 4      Female   67             14 Sales
## 4 4459-VLPQUH-3OL      4 3      Male     23             16 Sales
## 5 8158-SMTQFB-CNO      2 2      Male     26             16 Sales
```

Class Activity: select()

Which of the following can be used to:

- Q1. select all variables between CustomerID and Gender.
- Q2. select all variables other than those between CustomerID and Gender.
- Q3. select CustomerID and all variables that contain the word "Card".

filter(): filter observations based on values

- `filter()` identifies or selects observations in which a particular variable matches a specific value/condition.
- The condition(s) can be any kind of logical comparison and Boolean operators, such as:

<	Less than
>	Greater than
==	Equal to
<=	Less than or equal to
>=	Greater than or equal to
!=	Not equal to
%in%	Group membership
is.na	Is NA

!is.na	Is not NA
&,	Boolean AND, OR
xor	exactly or
!	not
any	any true
all	all true

Class Activity: filter()

Which of the following can be used to:

- Q4. filter for female customers only?
- Q5. filter for female customers that are greater than 45 years old AND live in region 3.
- Q6. filter for female customers that are greater than 45 years old OR live in region 3.

arrange(): order data by variables

- `arrange()` orders the data by variables in ascending (default) or descending order.
- For a descending order, use `desc()` within the `arrange()` function.

Class Activity: arrange()

Which of the following can be used to:

- Q7: select the variables CustomerID, Region, Gender, Age, HHIncome, Cardspend and save this as sub_cust.
- Q8: order sub_cust data by Age and CardSpendMonth (ascending order)
- Q9: order sub_cust data by Age (oldest to youngest) and CardSpendMonth (least to most)

mutate(): create new variables

- `mutate()` adds new variables while preserving the existing variables.
- `transmute()` creates a new variable and then drops the other variables.
- Here is the list of some useful functions used inside the `mutate()`.

<code>pmin(), pmax()</code>	Element wise min and max
<code>cummin(), cummax()</code>	Cumulative min and max
<code>cumsum(), cumprod()</code>	Cumulative sum and product
<code>between()</code>	Are values between a and b?
<code>cume_dist()</code>	Cumulative distribution of values
<code>cumall(), cumany()</code>	Cumulative all and any
<code>cummean()</code>	Cumulative mean

Class Activity: mutate()

Which of the following can be used to:

- Q10: create a ratio variable that computes the ratio of CardSpendMonth to HHIncome using sub_cust data
- Q11: create two variables: ratio1 = CardSpendMonth / HHIncome and ratio2 = CardSpendMonth / Age

summarise(): summarise data by functions of choice

- `summarise()` (or `summarize()`) performs the majority of summary statistics.

Functions	Usage
<code>min()</code> , <code>max()</code>	Minimum and maximum values
<code>mean()</code>	Mean value
<code>median()</code>	Median value
<code>sum()</code>	Sum of values
<code>var()</code> , <code>sd()</code>	Variance and standard deviation of a vector
<code>first()</code>	First value in a vector
<code>last()</code>	Last value in a vector
<code>nth()</code>	Nth value in a vector
<code>n()</code>	The number of values in a vector
<code>n_distinct()</code>	The number of distinct values in a vector

- All functions in this list takes a vector of values and returns a single summary value.

group_by() + summarise() function

- If we want to take the summary statistics grouped by a variable, then we need to use another function called `group_by()`.
- `group_by()` along with `summarise()` functions will allow us to take and compare summary statistics grouped by a variable.

Class Activity: `group_by()` + `summarise()`

Which of the following can be used to:

- Q12: compute the average `CardSpendMonth` across all customers in our `sub_cust` data.
- Q13: compute the average `CardSpendMonth` for each `Gender`.
- Q14: compute the average `CardSpendMonth` for each `Gender` and `Region`.

Joining data sets

- Often we have separate data frames that can have common and differing variables for similar observations (relational data sets).

```
band_members
```

```
## # A tibble: 3 x 2
##   name band
##   <chr> <chr>
## 1 Mick  Stones
## 2 John  Beatles
## 3 Paul  Beatles
```

```
band_instruments
```

```
## # A tibble: 3 x 2
##   name plays
##   <chr> <chr>
## 1 John  guitar
## 2 Paul  bass
## 3 Keith guitar
```

```
band_instruments2
```

```
## # A tibble: 3 x 2
##   artist plays
##   <chr> <chr>
## 1 John  guitar
## 2 Paul  bass
## 3 Keith guitar
```

Joining data sets

- `dplyr` offers three sets of joining functions to provide alternative ways to join data frames.
 - **Mutating joins**: add new variables to one data frame from matching observations in another.
 - **Filtering joins**: filter observations from one data frame based on whether or not they match an observation in the other table.
 - **Set operations**: treat observations as if they were set elements.
 - **Merging data sets**: merge data frames by row and column
- Please refer to the cheatsheets under the drive folder for details.

Mutating joins

- ``left_join()``: prioritizes left dataset
 - ``right_join()``: prioritizes right dataset
 - ``inner_join()``: only retains rows in both datasets
 - ``full_join()``: retains all rows
-
- Each mutating join takes an argument `by` that controls which variables are used to match observations in the two data sets.
 - `NULL`: The default value. `dplyr` will use all variables that appear in both tables, a natural join.
 - A character vector, `by = "x"`

Example: Mutating joins

```
library(dplyr)

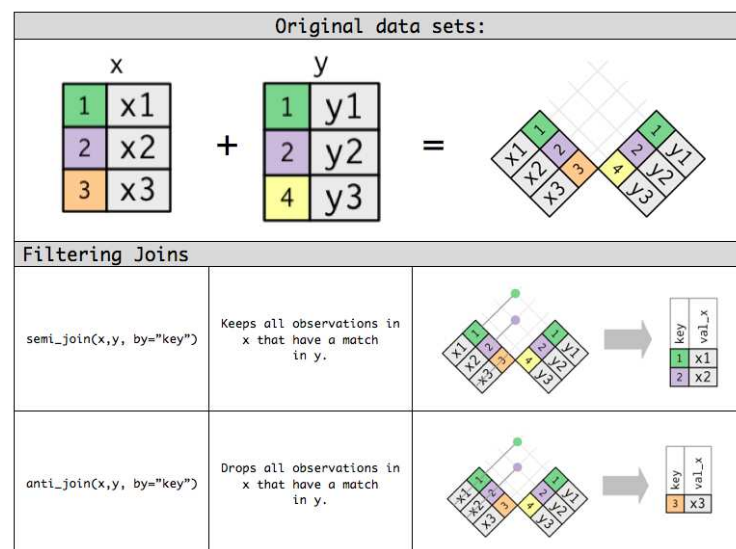
band_members
band_instruments
band_instruments2

# "Mutating" joins add variables to the LHS

band_members %>% inner_join(band_instruments)
band_members %>% left_join(band_instruments)
band_members %>% right_join(band_instruments)
band_members %>% full_join(band_instruments)
```

Filtering joins

- ``semi_join(x, y)``: keeps all observations in x that have a match in y.
- ``anti_join(x, y)``: drops all observations in x that have a match in y.



Adapted from Wickham, Hadley, and Garrett Golemund. 2016. R for Data Science: Import, Tidy, Transform, Visualize, and Model Data. O'Reilly Media, Inc.

Example: Filtering joins

```
# "Filtering" joins keep cases from the LHS  
band_members %>% semi_join(band_instruments)  
band_members %>% anti_join(band_instruments)
```

Set operations

- ``intersect(x, y)``: return only observations in both x and y.
- ``union(x, y)``: return unique observations in x and y.
- ``setdiff(x, y)``: return observations in x, but not in y.

Merging data sets

- ``bind_rows(x, y)``: Append y to x as new rows.
- ``bind_cols(x, y)``: Append y to x as new columns.

Original data sets:																					
y			z																		
x1	x2	+	x1	x2	=																
A	1		B	2																	
B	2		C	3																	
C	3		D	4																	
Merging data sets																					
bind_rows(y, z)	Append z to y as new rows			<table><tr><th>x1</th><th>x2</th><th>x1</th><th>x2</th></tr><tr><td>A</td><td>1</td><td>B</td><td>2</td></tr><tr><td>B</td><td>2</td><td>C</td><td>3</td></tr><tr><td>C</td><td>3</td><td>D</td><td>4</td></tr></table>		x1	x2	x1	x2	A	1	B	2	B	2	C	3	C	3	D	4
x1	x2	x1	x2																		
A	1	B	2																		
B	2	C	3																		
C	3	D	4																		
bind_cols(y, z)	Append z to y as new columns			<table><tr><th>x1</th><th>x2</th></tr><tr><td>A</td><td>1</td></tr><tr><td>B</td><td>2</td></tr><tr><td>C</td><td>3</td></tr><tr><td>B</td><td>2</td></tr><tr><td>C</td><td>3</td></tr><tr><td>D</td><td>4</td></tr></table>		x1	x2	A	1	B	2	C	3	B	2	C	3	D	4		
x1	x2																				
A	1																				
B	2																				
C	3																				
B	2																				
C	3																				
D	4																				

Adapted from Wickham, Hadley, and Garrett Grolemund. 2016. R for Data Science: Import, Tidy, Transform, Visualize, and Model Data. O'Reilly Media, Inc.

Functions to Remember for Week 5 (dplyr)

Operator/Function	Description
<code>filter</code>	pick observations based on their values
<code>></code> , <code>>=</code> , <code><</code> , <code><=</code> , <code>!=</code> , <code>==</code>	comparison operators
<code>arrange</code>	re-order rows
<code>desc</code>	order in descending order
<code>select</code>	select variables
<code>starts_with</code> , <code>ends_with</code> , <code>contains</code> , etc.	select variables based on patterns
<code>mutate</code> , <code>transmute</code>	create new variables
<code>summarise</code>	summarize data
<code>group_by</code>	group based on categorical variables
<code>%>%</code>	pipe operator to chain together functions

- Practice!

Class Worksheet



- Working in small groups, complete the following class worksheet

Week 5 Class Worksheet

- Once completed, feel free to work on your Assignment and/or Skill Builders

[Return to Data Preprocessing Website](#)